




(a) Reference-geometry rates**(b) Paired difference****(c) Discordant trials**

Equal cap is not equal computation: mean generation calls actually used (correction vs best-of-three) are TCN 2.72/2.99, QP 2.67/2.98, heuristic 1.84/1.81.

The uncorrected arms (cap 1) and the one-correction / best-of-two arms (cap 2) are not budget-matched and are not compared here.

36-scene cluster bootstrap, 30,000 resamples, descriptive. Reference geometry only: no tracked or obstacle-present outcome is claimed.